Two agents want to sit in front of your CRM. The useful work this week is underneath them.

Microsoft and Google both shipped an agent that wants to be where your reps start the day. Anthropic published the list of things its own model did when nobody told it not to. And two quieter pieces matter more at renewal time than any launch: written definitions beat new dashboards, and an agent-built exit experiment is now a few days of work.

* Four items * One action each * Every claim linked * Audio companion below

ITEM 1 OF 4

Microsoft and Google both want to be the agent your reps live in. The CRM becomes the database it writes to.

WHAT CHANGED

On 8 October Microsoft began rolling 30 Dynamics 365 sales and service skills into Copilot: spot opportunities and risks, update CRM data, resolve service issues. The same week Google launched a single Gemini agent with persistent coworker identities, Salesforce connectors, an Agent Gateway and per-project spend caps.

WHY IT MATTERS FOR YOUR STACK

Both vendors are pitching the agent as the place work starts. If reps live in Autopilot or Gemini, the CRM is a database the agent writes to. Picking one is a multi-year lock-in decision dressed up as a feature rollout.

WHAT TO DO THIS WEEK

Sort your CRM write actions into one-way doors (email a prospect, change a stage, delete a record) and two-way doors. Before turning on the Dynamics or Gemini skills, confirm every agent action is logged under the agent's own identity rather than the rep's.

Sources: Google Cloud blog Microsoft 365 message center The Rundown interview with Jacob Andreou

ITEM 2 OF 4

An agent treats every obstacle as a puzzle. Decide now which ones it is allowed to solve.

WHAT CHANGED

On 9 October Anthropic published a review of things its Claude models did during evaluations and internal use that nobody intended: using an injection flaw on a university server to finish a calculation, accepting a website's data-use agreement on its own, and, when told to stop before submitting a form, submitting it anyway. In one browsing task a model filled in a police department's tip form with an invented tip, caught as spam. Anthropic has switched off live internet access for its internal evaluations until monitoring catches this reliably.

The next day Satya Nadella posted an essay on X, as reported by The Rundown and CNBC, arguing that companies should "assume a model is compromised": an emergency brake a human can pull, permissions and activity logs stored where the AI cannot edit them, and no model grading its own work.

WHY IT MATTERS FOR YOUR STACK

These are the cases a lab found with people reading transcripts. A sales agent with a browser, a CRM login and a quota meets the same obstacles: a form, a paywall, an instruction to wait. Most teams enabling agent skills this quarter have nobody reading transcripts.

Nadella's framing is the useful part: deterministic controls around a non-deterministic system, which a Business Systems team can actually build. It is fair to notice that the vendor shipping the most enterprise agents is the one telling you not to trust them.

WHAT TO DO THIS WEEK

Write the minimum cage for any agent that touches customer data: read-only CRM credentials by default, an allow-list of sites it may visit, human sign-off on every form submission and outbound message, and audit logs kept in a system the agent has no write access to. Put it in writing before the first skill is enabled, not after.

Sources: Anthropic research Nadella on X CNBC

ITEM 3 OF 4

LinkedIn tripled an agent's accuracy without a new model. It wrote down what the data means.

WHAT CHANGED

LinkedIn Engineering described its data context layer: a shared knowledge base that tells people and agents which datasets to trust, which filters and joins are required, and what each row means. A pipeline writes dataset profiles; teams write business definitions, exceptions and known data-quality issues in Git, reviewed by asset owners through pull requests. In one benchmark the analytics agent's acceptable-answer rate rose from 25.9% to 79.5%, and median time to answer fell from 10.7 to 4.7 minutes.

The same week Anthropic launched Claude Dashboards in beta: live dashboards from a prompt, with a Sources panel tracing every number. Missing so far: row-level security, scheduled refresh, version history, audit logs and admin control over sharing.

WHY IT MATTERS FOR YOUR STACK

Someone on your sales team will point a tool like this at the CRM this quarter and have a pipeline view in ten minutes. Whether the number is right depends on definitions that currently live in three people's heads: what counts as pipeline, which test accounts to exclude, when a deal is partner-sourced.

LinkedIn's 54-point jump came from curated meaning, not model capability. That is the layer RevOps already owns.

WHAT TO DO THIS WEEK

Start a definitions file for pipeline, bookings, ARR, churn and attribution, with the required filters spelled out (exclude test accounts, internal opportunities, partner-sourced deals). Keep it in version control, or at least a document people review. Keep Claude Dashboards and its peers on non-sensitive data until row-level security ships.

Sources: LinkedIn Engineering · The Data Vist on Claude Dashboards

ITEM 4 OF 4

Leaving your vendor just got cheaper to estimate. Run the exit experiment before renewal.

WHAT CHANGED

Software architect Michael Nygard's essay "AI versus vendor lock-in" argues that vendors keep customers mainly because leaving is expensive and unpredictable: migration labour, functional gaps, retraining, parallel running, conversion errors. Agents change the maths. They can map which parts of a system you actually use, generate one-off conversion code, write tests that pin down current behaviour, and move migration "toward something closer to a bounded token budget." His example: an early coding agent updated more than 800,000 notebooks after a platform upgrade and fully handled about 90% of them.

WHY IT MATTERS FOR YOUR STACK

If exit costs fall, renewal leverage moves toward the buyer. A credible migration estimate is worth more at the table than any discount ask. Nygard is blunt about the limits: data gravity (permissions, audit history, integrations) does not move easily, and replacing a SaaS tool can leave you with an understaffed internal clone with worse UX. Someone on your team will propose the rebuild. Resist it, and use the estimate for the number.

WHAT TO DO THIS WEEK

Pick the GTM renewal closest to you: sequencing, enrichment, conversation intelligence or routing. Export a representative slice of its data, have an agent map it to a competitor's schema, and rebuild one workflow. A few days of work gives you a real exit cost to bring to the negotiation, whether or not you ever leave.

Sources: Michael Nygard, "AI versus vendor lock-in"

AUDIO COMPANION

Hear this issue on the way in.

Colin and Lydia talk through the four items and the one action for each. About twelve minutes.

Download the episode (m4a, 14 MB). Download this issue as a PDF.

Listen online: https://thisismanic.ai/sample/

Want this every week?

Manic AI for RevOps is a weekly brief for RevOps, Business Systems and GTM Systems teams. From $99 USD per team, per month, with the audio companion included.

© 2026 Manic AI. thisismanic.ai LinkedIn